Comparative Analysis of Vectorization Models and Retrieval Methods for Finding Similar Scientific Papers

Arman Tursynbayev
School of Computer Science
Astana IT University
Astana, Kazakhstan
www.armani7775@gmail.com

Zhanna Sadirmekova
School of Computer Science
Astana IT University
Astana, Kazakhstan
Zhanna.Sadirmekova1@astanait.edu.kz

***Abstract*—Literature search often starts from one paper: the reader has a relevant article and wants others on the same problem. We compare widely used text representations for this query-by-document task on arXiv computer-science preprints: BM25, TF-IDF, latent semantic analysis (LSA), three pretrained dense encoders (all-MiniLM-L6-v2, SPECTER2 and BGE-small) and a BM25–dense hybrid. Two corpora of 10,000 papers were collected through the arXiv API, one covering three weeks of May 2026 and one sampled across 2020–2024. Every eligible paper serves as a query with itself removed from the candidates; 1,000 queries per corpus are reserved for parameter selection and the remaining 8,368 and 8,497 for testing. A retrieved paper counts as relevant when it shares the query's primary arXiv category, a label that no compared system computes. All three dense encoders beat every sparse baseline on both corpora, also after the baselines are tuned: on the May corpus P@5 is 0.520 for SPECTER2, 0.512 for BGE-small and 0.510 for MiniLM against 0.456 for BM25, and the differences are significant under paired randomization tests with Holm correction. A tuned BM25 component adds little to a dense score. A blind problem-level reading of 10 queries by the first author tells a different story: same-problem precision is low for every system (0.10–0.20) and the dense encoders are not ahead, suggesting that category precision reflects topical closeness more than usefulness to a reader.**

***Index Terms*—similar-paper search, query by document, BM25, latent semantic analysis, sentence embeddings, SPECTER2, hybrid retrieval, arXiv**

I. Introduction

A common way to explore a research topic is to start from one paper and ask for more like it. Digital libraries expose this as "related articles" or "similar papers," and intelligent services that organize access to scientific and educational resources rely on the same operation [1], [2]. The engineering choice behind it is usually between a lexical ranker such as BM25 [3], a vector-space or latent-semantic representation [4], [5], and a pretrained dense encoder [6], [7]. Published comparisons do not settle the choice. On heterogeneous retrieval benchmarks BM25 and dense models trade the lead from task to task [8], [9], and on large citation-recommendation pools BM25 remained competitive with transformer article encoders [10].

Evaluating similar-paper search also has a labeling problem. Expert relevance judgments for paper-to-paper similarity are expensive and exist for only tens of queries; CSFCube, for example, has 50 [11]. Two cheaper shortcuts are common and both bias the result. If "relevant" is defined as "close under an embedding," that embedding wins by construction. If the source paper stays in the index and the query is a fragment of its title, every ranker retrieves the source paper at rank one and reciprocal rank says little about the rest of the list. We avoid both: the query paper is removed from its own candidate list, and relevance is the primary arXiv category, which authors and moderators assign and which none of the compared systems computes.

This paper is an empirical comparison, not a new ranking model. Its contributions are:

- a controlled comparison of seven retrieval configurations on two arXiv corpora, a recent three-week slice and a five-year sample, using every eligible paper as a query rather than a small sample;
- a sensitivity and ablation analysis covering BM25 k_1 and b, TF-IDF weighting, LSA rank, the hybrid weight and score normalization, and the text fields, with all choices made on a held-out query set;
- paired significance tests with bootstrap confidence intervals for every headline difference;
- a blind, problem-level reading of pooled results that measures how often the category proxy and same-problem relevance disagree.

Code, query lists, per-query scores and the manual labels are publicly released (see Data and Code Availability).

II. Related Work

Paper recommendation and scientific document representations. Surveys of paper recommender systems group them into content-based, collaborative and graph-based families and note that evaluation protocols differ widely between studies [12], [13]. Among content-based encoders, SPECTER trains on citation triplets [7], SciNCL samples training pairs from

citation-graph neighborhoods [14], and SPECTER2 adds task-format adapters after SciRepEval showed that one embedding is a poor compromise across scientific tasks [15]. ASPIRE matches papers through several aspect vectors [16], and CSFCube provides faceted expert judgments for query-by-document search in computer science [11]. LitSearch, a recent benchmark of literature-search questions, reports a large recall gap between BM25 and modern dense retrievers [17]; its queries are natural-language questions, whereas ours are whole papers.

Lexical versus dense retrieval. BEIR showed that zero-shot dense retrievers do not uniformly beat BM25 [8]. MTEB and its multilingual extension MMTEB find that no single embedding model leads on all tasks [9], [18]. General-purpose encoders such as BGE are trained on large weakly supervised pair collections [19], while compact models such as MiniLM are distilled from larger transformers [20] and fine-tuned as sentence encoders [6].

Hybrid retrieval. Sparse and dense scores are often combined because they miss different documents [21]. Bruch et al. show that a tuned convex combination of normalized scores usually outperforms reciprocal rank fusion (RRF) and that the result is fairly insensitive to the normalization choice [22]; RRF itself combines ranks rather than scores [23]. We test both, with the weight chosen on held-out queries.

III. Data

Both corpora were downloaded with the `arxiv` Python client from the public arXiv API using the query `cat:cs.AI OR cat:cs.LG OR cat:cs.CL OR cat:cs.IR`. Each record keeps the identifier, title, abstract, author list, category list, primary category, and first-version and latest-version dates. The category list stores the `cs.*` tags of a paper; the primary category is stored even when it lies outside computer science, so cross-listed papers such as `stat.ML` or `eess.IV` keep their own primary label.

May 2026 corpus. This collection was also used in our earlier 100-query evaluation: the 10,000 most recent matching records returned on 31 May 2026, sorted by submission date, with publication dates from 5 to 28 May 2026. The API stopped returning pages at 10,000 records.

2020–2024 corpus. To test whether the results depend on a three-week window, we built a second corpus. Each month of 2020–2024 was split into four windows starting on days 1, 8, 15 and 22, and the first 45 papers by submission date were taken from each window. Of the 240 windows, the two covering 1–14 January 2024 failed with HTTP 429 rate-limit responses and were not retried, so that the stored sample is not changed after it was drawn. After removing 90 duplicates, 10,710 candidates remained and 10,000 were kept by a seeded uniform draw.

Table I summarizes both files. Neither contains a paper twice under different version suffixes, every record has a publication date inside its stated range, and every abstract consists of at least 90% ASCII letters, so no language filter was needed. Two kinds of near-duplicate hits could inflate precision: papers by the same authors and near-identical texts. Across all systems, 1.5–2.9% of the 41,840 top-five slots of the May test queries share at least one author name with the query (1.4–2.8% for 2020–2024), and 34–46 slots per system have a title-word Jaccard of 0.5 or more with the query. Both effects are small: removing every candidate that shares an author with the query changes P@5 by at most 0.008 (0.005 for 2020–2024) and every dense encoder stays ahead of every sparse baseline.

TABLE I
THE TWO CORPORA.

	May 2026	2020–2024
Records	10,000	10,000
Publication dates	2026-05-05 – 05-28	2020-01 – 2024-12
Distinct primary categories	107	104
Primary in query categories	69.2%	59.7%
cs.LG / cs.CL	34.6% / 15.7%	32.1% / 17.9%
cs.AI / cs.CV	16.7% / 9.3%	6.2% / 11.4%
Categories per paper (mean)	1.79	1.87
Abstract words (median)	194	175
Revised after v1	11.3%	42.0%
Queries: eligible/dev/test	9,368/1,000/8,368	9,497/1,000/8,497
Chance precision (test)	0.196	0.165

Queries and splits. A paper is eligible as a query when its primary category has at least 30 papers in the corpus, so that five same-category hits remain possible. From the eligible papers, 1,000 are drawn with seed 43 as a development set; they are used only to choose parameters. All other eligible papers form the test set, and every number in the main tables comes from it.

IV. Retrieval Methods

All systems index the concatenation of title and abstract. For a query paper q the candidates are $\mathcal{C}_q = \mathcal{D} \setminus \{q\}$, and each system returns the top 100 candidates.

BM25 [3] scores a candidate d as

$$s(q,d) = \sum_{t \in q} \mathrm{idf}(t)\, \frac{f_{t,d}\,(k_1+1)}{f_{t,d} + k_1\Big(1 - b + b\,|d|/\overline{|d|}\Big)}, \quad (1)$$

with $\mathrm{idf}(t) = \ln\Big(1 + \frac{N - n_t + 0.5}{n_t + 0.5}\Big)$, where $f_{t,d}$ is the frequency of term t in d, n_t the number of documents containing t, and N the corpus size. Each distinct query term is counted once. Tokens are lowercase alphanumeric strings longer than two characters, excluding the scikit-learn English stop list and pure numbers. The defaults $k_1 = 1.5$ and $b = 0.75$ are those of the `rank_bm25` package; Pyserini uses $k_1 = 0.9$ and $b = 0.4$ [24], a setting included in our sensitivity grid.

TF-IDF [4] uses scikit-learn with the same stop list, terms in at least 3 and at most 90% of documents, and at most 20,000 terms; vectors are L2-normalized and compared by cosine similarity.

LSA [5] applies truncated SVD, $X \approx U_k \Sigma_k V_k^{\top}$, to the TF-IDF matrix and represents each paper by its L2-normalized row of $U_k \Sigma_k$; the default rank is $k = 128$.

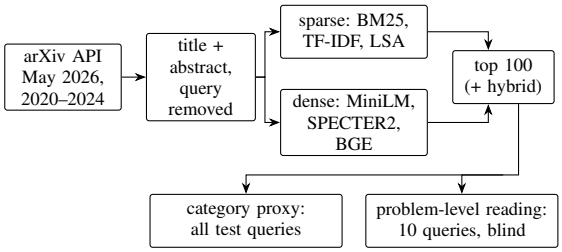


Fig. 1. Evaluation pipeline. The query paper is removed from its own candidate list before ranking.

Dense encoders. We use three public checkpoints without fine-tuning: `all-MiniLM-L6-v2` (384 dimensions, input truncated at 256 word pieces) [6], [20]; SPECTER2 with its proximity adapter (768 dimensions, title and abstract joined by the separator token, 512 word pieces, CLS pooling) [15]; and `bge-small-en-v1.5` (384 dimensions, 512 word pieces) [19]. Vectors are L2-normalized and ranked by cosine similarity.

Hybrid. The hybrid score combines a dense score and BM25:

$$h(q,d) = w\,\hat{s}_{\text{dense}}(q,d) + (1-w)\,\hat{s}_{\text{BM25}}(q,d), \qquad (2)$$

where $\hat{s}$ is min-max normalized over $\mathcal{C}_q$. The pre-specified configuration is MiniLM with $w = 0.6$, the value fixed in the configuration of our earlier prototype. In the sensitivity analysis we also vary w from 0 to 1, replace min-max by z-score normalization, and test RRF with constant 60 [23]. Normalization excludes the query paper itself; including it would compress the BM25 component, because a paper's BM25 score against itself is far larger than against any other paper.

Fig. 1 summarizes the pipeline.

V. Evaluation Protocol

Relevance proxy. A candidate is relevant when its primary category equals the query's. We use this label because it is available for every paper, so all test queries can be scored, and because no compared system produces it. It is an imperfect proxy for "the same research problem": a paper on the same problem can be filed under another category, and two papers in one category can address different problems. The problem-level reading below measures both failure modes directly.

Metrics. Precision at k is the fraction of the top k that is relevant; we report $k = 5$ and $k = 10$. Mean reciprocal rank (MRR) uses the rank of the first relevant candidate within the top 100. For nDCG@10 the gain is graded, $g(q,d) = |C_q \cap C_d| / |C_q \cup C_d|$, the Jaccard overlap of the two category lists, and

$$\text{nDCG@10} = \frac{1}{Z_q} \sum_{i=1}^{10} \frac{2^{g(q,d_i)} - 1}{\log_2(i+1)}, \qquad (3)$$

where Z_q is the same sum for the ten candidates with the largest gain. A random ranking would reach an expected precision of $(n_c - 1)/(N - 1)$ for a query in category c of size n_c; averaged over the test queries this chance level is 0.196 for May 2026 and 0.165 for 2020–2024.

Statistics. For each pair of systems we compute per-query differences, a 95% bootstrap interval of the mean difference (10,000 resamples), and a two-sided paired randomization test with 10,000 sign flips, which is well suited to retrieval metrics [25]. Within each metric, p-values are adjusted with the Holm procedure [26]. With 10,000 permutations the smallest attainable raw p is about 10^{-4}, so after the Holm adjustment over 15 comparisons the floor is about 0.0015, reported as $p < 0.002$.

Problem-level reading. Ten test queries from the May corpus were drawn, one from each of the ten most frequent primary categories in the test set (seed 11). The top five of all seven systems were pooled, giving 212 distinct candidates. The first author read the title and abstract of each query and candidate and marked a candidate relevant if it studies the same research problem: the same task applied to the same kind of object, a paper one would expect to cite in the related-work section. A shared field or technical term was not sufficient. The pool was shuffled, and neither the retrieving system nor any arXiv category was shown. Reader P@5 is the share of a system's own top five marked relevant; every top-five candidate is in the pool, so no hit is unjudged. This is a single-annotator check of the proxy, not a test collection.

TABLE II
Test-set results with pre-specified parameters. Relevance: same primary category; nDCG@10 uses graded category overlap. Every 95% interval for P@5 and nDCG@10 has a half-width of at most 0.007. Best value per column in bold.

	May 2026			2020–2024		
Method	P@5	nDCG@10	MRR	P@5	nDCG@10	MRR
BM25	0.456	0.440	0.647	0.481	0.460	0.662
TF-IDF	0.429	0.418	0.618	0.452	0.434	0.623
LSA	0.454	0.437	0.627	0.467	0.458	0.633
MiniLM	0.510	0.467	0.679	0.561	0.509	0.708
SPECTER2	**0.520**	**0.474**	**0.685**	**0.567**	**0.515**	**0.711**
BGE-small	0.512	0.466	0.684	0.542	0.506	0.695
Hybrid (0.6)	0.509	0.469	0.684	0.548	0.503	0.705

VI. Results

A. Main comparison

Table II gives the test-set results with pre-specified parameters. On both corpora each of the three dense encoders is ahead of each sparse baseline, and every system is far above the chance level. SPECTER2, the only encoder trained on scientific citations, is first on both corpora, but its lead over MiniLM is small (+0.010 and +0.006 P@5). BGE-small matches MiniLM on the May corpus and falls behind it on the 2020–2024 corpus. Among the sparse systems BM25 is best on both corpora; LSA is close to BM25 and clearly ahead of the TF-IDF matrix it is computed from.

The paired tests in Table III show that the advantage of each dense encoder over the best sparse baseline is small

in absolute terms but precisely estimated: with more than 8,000 paired queries the bootstrap intervals exclude zero and the Holm-adjusted p-values are at the permutation floor. The pre-specified MiniLM hybrid ($w = 0.6$) is indistinguishable from MiniLM on the May corpus (-0.001, $p_{\text{Holm}} = 0.95$) and slightly worse on the 2020–2024 corpus (-0.013, $p_{\text{Holm}} < 0.002$). BM25 is ahead of TF-IDF on both corpora, a difference that our earlier 100-query evaluation could not resolve; that sample also underestimated BM25, whose P@5 on it was 0.386.

The result lists overlap little. On the May corpus the mean Jaccard overlap of the top 10 is 0.13 between MiniLM and BM25, 0.12 between MiniLM and LSA, and only 0.20–0.26 among the three dense encoders, so systems with similar precision still return largely different papers.

B. Per-category results

Table IV splits P@5 for the four largest query categories of the May corpus. The gap between sparse and dense systems is smallest in `cs.LG` and `cs.AI`, where precision is bounded by heavy cross-listing, and largest in `cs.CL` and `cs.CV`, where task names and model names recur across categories and term matching alone attracts off-category papers.

C. Sensitivity, ablation and hybrid weights

Table V reports development-set results for the parameter grids and the configuration selected by nDCG@10, scored once on the test set. Tuning moves the sparse baselines by at most 0.033 in test P@5 on the May corpus and 0.039 on the 2020–2024 corpus, where the configuration is selected on that corpus's own development queries; on both corpora no tuned sparse configuration reaches the weakest dense encoder (at best 0.491 against 0.542 on 2020–2024). TF-IDF benefits most, from sublinear term frequency. For LSA, lower ranks are slightly better on this task, so the 128-dimensional default is not handicapped by too few components. The field ablation shows that the abstract carries most of the signal: title-only indexing loses 0.057–0.125 in dev P@5 for the sparse systems and 0.057 for MiniLM. Even from titles alone MiniLM (0.468) is above every sparse system indexed on title and abstract (at most 0.459).

Fig. 2 shows nDCG@10 of the MiniLM–BM25 hybrid across weights on the test sets. The curves rise from BM25 alone to a broad plateau: between $w = 0.6$ and $w = 0.9$ nDCG@10 varies by 0.003 on the May corpus and 0.008 on the 2020–2024 corpus. Z-score and min-max normalization give nearly the same curve, in line with [22], and RRF (0.471 and 0.504) is close to the best convex weight (0.472 and 0.510). The weight selected on development queries ($w = 0.8$ and 0.9) puts most of the mass on the dense score, and the tuned hybrid reaches MiniLM's precision but does not exceed the dense encoders by a meaningful margin (Table V).

D. Problem-level reading

Table VI compares reader P@5 with category P@5 on the ten judged queries. Of the 212 pooled candidates the first

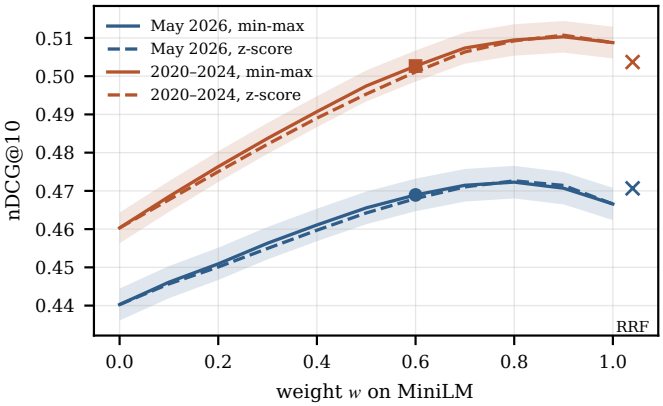


Fig. 2. nDCG@10 of $w\,\hat{s}_{\text{MiniLM}} + (1 - w)\,\hat{s}_{\text{BM25}}$ on the test sets ($w = 0$: BM25 only, $w = 1$: MiniLM only). Solid: min-max, dashed: z-score normalization; bands: 95% intervals; dots: the pre-specified $w = 0.6$; crosses: RRF.

author marked 18 as the same research problem; for 2 of the ten queries no pooled candidate qualified. Reader P@5 is low for every system, between 0.10 and 0.20, and it does not follow the ranking by category: BM25, TF-IDF and the hybrid reach 0.18–0.20, while LSA, MiniLM and SPECTER2 reach 0.10. Of the 18 same-problem papers, 9 were retrieved by none of the three dense encoders and 1 only by them. Category agreement over-counts far more often than it under-counts: of the 150 top-five slots of the dense encoders, 79 share the query's category without addressing the same problem, while 13 address the same problem from another category. With ten queries and one annotator these figures are descriptive, but they indicate that the dense advantage in category precision reflects topical closeness more than same-problem relevance.

VII. Discussion

The two corpora differ in time span, topic mix and the share of cross-listed papers, yet the ordering of systems is the same on both: the dense encoders first, BM25 and LSA next, TF-IDF last. The May corpus is therefore not an outlier caused by a three-week window. At the same time the margins are moderate. The best dense encoder improves category P@5 over the best sparse baseline by 6.4 points on the May corpus and 8.6 points on the 2020–2024 corpus, and the per-category split shows that the gain is concentrated in categories whose vocabulary is shared with other fields. A domain-specific encoder helps only slightly: SPECTER2 is ahead of the general-purpose MiniLM by about one point.

The problem-level reading qualifies the main tables. Category agreement counts same-problem papers filed elsewhere as misses and counts different-problem papers in the same category as hits, and on the judged queries the second error is much more frequent. The dense encoders are good at returning papers from the same area, which is what the category proxy rewards, but on these queries a paper about the same specific problem was more often found by term matching, plausibly because such papers share specific task,

TABLE III
PAIRED DIFFERENCES IN TEST P@5: MEAN DIFFERENCE, 95% BOOTSTRAP INTERVAL AND HOLM-ADJUSTED RANDOMIZATION p-VALUE.

	May 2026		2020–2024	
Comparison	Δ [95% CI]	p	Δ [95% CI]	p
MiniLM – BM25	+0.054 [+0.048, +0.060]	<0.002	+0.080 [+0.074, +0.085]	<0.002
SPECTER2 – BM25	+0.064 [+0.058, +0.070]	<0.002	+0.086 [+0.080, +0.092]	<0.002
BGE-small – BM25	+0.056 [+0.050, +0.062]	<0.002	+0.061 [+0.055, +0.067]	<0.002
SPECTER2 – MiniLM	+0.010 [+0.004, +0.015]	0.004	+0.006 [+0.001, +0.011]	0.014
BGE-small – MiniLM	+0.002 [−0.003, +0.007]	0.946	−0.019 [−0.023, −0.014]	<0.002
BGE-small – SPECTER2	−0.008 [−0.013, −0.002]	0.016	−0.025 [−0.030, −0.020]	<0.002
Hybrid (0.6) – MiniLM	−0.001 [−0.006, +0.004]	0.946	−0.013 [−0.017, −0.009]	<0.002
TF-IDF – BM25	−0.027 [−0.033, −0.021]	<0.002	−0.029 [−0.035, −0.023]	<0.002
LSA – TF-IDF	+0.024 [+0.018, +0.030]	<0.002	+0.015 [+0.009, +0.021]	<0.002

TABLE IV
TEST P@5 BY PRIMARY CATEGORY OF THE QUERY, MAY 2026 CORPUS. CHANCE IS $(n_c - 1)/(N - 1)$.

Cat.	n	Chance	BM25	TF-IDF	LSA	MiniLM	SPEC2	BGE
cs.LG	3,078	0.346	0.577	0.560	0.591	0.608	**0.627**	0.616
cs.AI	1,509	0.167	0.327	0.314	0.328	0.362	0.378	**0.386**
cs.CL	1,405	0.157	0.441	0.428	0.442	**0.540**	0.527	0.532
cs.CV	821	0.093	0.518	0.444	0.484	**0.648**	0.637	0.586

dataset or model names with the query. Category precision is useful for comparing systems on thousands of queries, because it is independent of the systems and cheap, but it should not be read as the share of results a researcher would keep, and it may overstate the benefit of dense retrieval for problem-level search. Ten queries are too few to rank systems by reader precision; the reading shows the size of the gap between the two notions of relevance, not a reversed ranking.

Limitations. The relevance label is a proxy, and the problem-level check covers 10 queries and one annotator, so it estimates the direction and rough size of the disagreement rather than precise reader precision. Both corpora come from four arXiv computer-science categories and contain titles and abstracts only, so the results do not cover full text, citation graphs or other disciplines. The dense encoders were used off the shelf; fine-tuning or larger models may change the margins. Two weekly windows of January 2024 are missing from the multi-year corpus because of API rate limits.

VIII. CONCLUSION

On two arXiv computer-science corpora, with every eligible paper used as a query and the query removed from its own results, pretrained dense encoders retrieve more same-category papers than BM25, TF-IDF and LSA. The advantage holds over three weeks of 2026 and over five years of 2020–2024, survives tuning of the sparse baselines on held-out queries, and is statistically reliable, although moderate in size. Mixing BM25 into a dense score with a tuned weight neither helps nor hurts category precision; a fixed weight of 0.6 can hurt it slightly. A blind problem-level reading of ten queries finds low same-problem precision for all systems and no dense advantage, which argues for reporting category precision together with problem-level judgments and for keeping a lexical component in similar-paper search. Future work will add expert judgments from several annotators and full-text and citation signals.

DATA AND CODE AVAILABILITY

The collection and evaluation scripts, the arXiv identifiers of both corpora, the development and test query lists, per-query scores, rankings and the manual labels are available at https://github.com/Armani7777/arxiv-similar-paper-retrieval.

ACKNOWLEDGMENT

The authors thank arXiv for open access to its metadata. Generative AI (Claude, Anthropic) was used to help write the analysis scripts and edit the manuscript text. The authors designed the study, ran all experiments, produced the manual relevance labels, verified all code, results and references, and take full responsibility for the content.

REFERENCES

[1] A. G. Batyrkhanov, Zh. B. Sadirmekova, M. A. Sambetbayeva, A. N. Nurgulzhanova, Zh. S. Ismagulova, and A. S. Yerimbetova, "Development of methods and technologies for creating intelligent scientific and educational internet resources," *Bull. Electr. Eng. Inform.*, vol. 11, no. 5, pp. 2968–2977, 2022, doi: 10.11591/eei.v11i5.3075.

[2] S. Serikbayeva, J. Tussupov, M. Sambetbayeva, A. Yerimbetova, Zh. Sadirmekova, A. Tungatarova, A. Batyrkhanov, and A. Zakirova, "Development of a model and technology of access to documents in scientific and educational activities," *East.-Eur. J. Enterp. Technol.*, vol. 6, no. 2(114), pp. 44–58, 2021, doi: 10.15587/1729-4061.2021.248506.

[3] S. Robertson and H. Zaragoza, "The probabilistic relevance framework: BM25 and beyond," *Found. Trends Inf. Retr.*, vol. 3, no. 4, pp. 333–389, 2009.

[4] G. Salton and C. Buckley, "Term-weighting approaches in automatic text retrieval," *Inf. Process. Manage.*, vol. 24, no. 5, pp. 513–523, 1988.

[5] S. Deerwester, S. T. Dumais, G. W. Furnas, T. K. Landauer, and R. Harshman, "Indexing by latent semantic analysis," *J. Amer. Soc. Inf. Sci.*, vol. 41, no. 6, pp. 391–407, 1990.

[6] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *Proc. EMNLP-IJCNLP*, 2019, pp. 3982–3992.

[7] A. Cohan, S. Feldman, I. Beltagy, D. Downey, and D. S. Weld, "SPECTER: Document-level representation learning using citation-informed transformers," in *Proc. ACL*, 2020, pp. 2270–2282.

[8] N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych, "BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2021.

[9] N. Muennighoff, N. Tazi, L. Magne, and N. Reimers, "MTEB: Massive text embedding benchmark," in *Proc. EACL*, 2023, pp. 2014–2037.

TABLE V

SENSITIVITY ON THE 1,000 DEVELOPMENT QUERIES (MAY 2026). RANGE: DEV NDCG@10 ACROSS THE GRID. THE CONFIGURATION WITH THE BEST DEV NDCG@10 IS THEN SCORED ONCE ON THE TEST SET (P@5, DEFAULT → SELECTED). BOTTOM: DEV P@5 BY INDEXED FIELD.

System	Grid	Dev range	Selected	Test P@5
BM25	$k_1 \in \{0.9, 1.2, 1.5, 2\}$, $b \in \{.4, .6, .75, .9\}$	0.429–0.437	k1=2.0, b=0.9	0.456 → 0.460
TF-IDF	min_df$\in \{1, 3, 5\}$, sublinear tf	0.412–0.440	min_df=1, sublin.	0.429 → 0.462
LSA	$k \in \{64, 128, 256, 512\}$	0.419–0.435	dim=64	0.454 → 0.458
MiniLM+BM25	$w \in [0, 1]$, min-max/z, RRF	0.434–0.471	z-score, w=0.8	0.510† → 0.516
SPECTER2+BM25	$w \in [0, 1]$, min-max/z, RRF	0.434–0.477	min-max, w=0.9	0.520† → 0.524
BGE-small+BM25	$w \in [0, 1]$, min-max/z, RRF	0.434–0.465	min-max, w=0.8	0.512† → 0.518
		Title	Abstract	Title+abstract
BM25	dev P@5	0.387	0.449	0.454
TF-IDF	dev P@5	0.378	0.430	0.435
LSA	dev P@5	0.334	0.446	0.459
MiniLM	dev P@5	0.468	0.521	0.525

† dense model alone.

TABLE VI

PROBLEM-LEVEL READING OF 10 MAY-2026 TEST QUERIES (ONE ANNOTATOR, BLIND TO SYSTEM AND CATEGORY). READER P@5: SAME RESEARCH PROBLEM. CATEGORY P@5: SAME PRIMARY CATEGORY, SAME LISTS. RIGHT: THE 50 TOP-FIVE SLOTS OF EACH SYSTEM SPLIT BY THE TWO LABELS.

	Reader	Category	Slots (problem / category)			
Method	P@5	P@5	yes/yes	yes/no	no/yes	no/no
BM25	0.18	0.40	3	6	17	24
TF-IDF	0.20	0.36	6	4	12	28
LSA	0.10	0.44	2	3	20	25
MiniLM	0.10	0.62	1	4	30	15
SPECTER2	0.10	0.52	1	4	25	20
BGE-small	0.16	0.54	3	5	24	18
Hybrid (0.6)	0.20	0.56	3	7	25	15

[10] Z. Medić and J. Šnajder, "Large-scale evaluation of transformer-based article encoders on the task of citation recommendation," in *Proc. 3rd Workshop Scholarly Document Processing*, 2022, pp. 19–31.

[11] S. Mysore, T. O'Gorman, A. McCallum, and H. Zamani, "CSFCube: A test collection of computer science research articles for faceted query by example," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2021.

[12] C. K. Kreutz and R. Schenkel, "Scientific paper recommendation systems: A literature review of recent publications," *Int. J. Digit. Libr.*, vol. 23, no. 4, pp. 335–369, 2022.

[13] I. Pinedo, M. Larrañaga, and A. Arruarte, "Recent advances and trends in research paper recommender systems: A comprehensive survey," arXiv:2508.08828, 2025.

[14] M. Ostendorff, N. Rethmeier, I. Augenstein, B. Gipp, and G. Rehm, "Neighborhood contrastive learning for scientific document representations with citation embeddings," in *Proc. EMNLP*, 2022, pp. 11670–11688.

[15] A. Singh, M. D'Arcy, A. Cohan, D. Downey, and S. Feldman, "SciRepEval: A multi-format benchmark for scientific document representations," in *Proc. EMNLP*, 2023, pp. 5548–5566.

[16] S. Mysore, A. Cohan, and T. Hope, "Multi-vector models with textual guidance for fine-grained scientific document similarity," in *Proc. NAACL*, 2022, pp. 4453–4470.

[17] A. Ajith, M. Xia, A. Chevalier, T. Goyal, D. Chen, and T. Gao, "LitSearch: A retrieval benchmark for scientific literature search," in *Proc. EMNLP*, 2024, pp. 15068–15083.

[18] K. Enevoldsen *et al.*, "MMTEB: Massive multilingual text embedding benchmark," in *Proc. ICLR*, 2025.

[19] S. Xiao, Z. Liu, P. Zhang, N. Muennighoff, D. Lian, and J.-Y. Nie, "C-Pack: Packed resources for general Chinese embeddings," in *Proc. 47th Int. ACM SIGIR Conf. Res. Develop. Inf. Retr.*, 2024.

[20] W. Wang, F. Wei, L. Dong, H. Bao, N. Yang, and M. Zhou, "MiniLM: Deep self-attention distillation for task-agnostic compression of pre-trained transformers," in *Proc. NeurIPS*, 2020, pp. 5776–5788.

[21] Y. Luan, J. Eisenstein, K. Toutanova, and M. Collins, "Sparse, dense, and attentional representations for text retrieval," *Trans. Assoc. Comput. Linguist.*, vol. 9, pp. 329–345, 2021.

[22] S. Bruch, S. Gai, and A. Ingber, "An analysis of fusion functions for hybrid retrieval," *ACM Trans. Inf. Syst.*, vol. 42, no. 1, 2024.

[23] G. V. Cormack, C. L. A. Clarke, and S. Büttcher, "Reciprocal rank fusion outperforms Condorcet and individual rank learning methods," in *Proc. 32nd Int. ACM SIGIR Conf. Res. Develop. Inf. Retr.*, 2009, pp. 758–759.

[24] J. Lin, X. Ma, S.-C. Lin, J.-H. Yang, R. Pradeep, and R. Nogueira, "Pyserini: A Python toolkit for reproducible information retrieval research with sparse and dense representations," in *Proc. 44th Int. ACM SIGIR Conf. Res. Develop. Inf. Retr.*, 2021, pp. 2356–2362.

[25] M. D. Smucker, J. Allan, and B. Carterette, "A comparison of statistical significance tests for information retrieval evaluation," in *Proc. 16th ACM Conf. Inf. Knowl. Manage. (CIKM)*, 2007, pp. 623–632.

[26] S. Holm, "A simple sequentially rejective multiple test procedure," *Scand. J. Statist.*, vol. 6, no. 2, pp. 65–70, 1979.